

# ECE 60827: Lab 1 Report

## SAXPY Execution:

The SAXPY kernel profiling data is shown below for three different inputs of VectorSize (see Table 1). The inputs all vary dramatically ranging from 1024 to 131072.

Input Size	Kernel / API Call	SAXPY Program			
		Execution Time (ms)	Percentage of Total Time (%)		
1024	cudaMemcpy HtoD	2.59E-03	38.94		GPU
	cudaMemcpy DtoH	2.02E-03	30.77		API
	saxpy_gpu()	2.05E-03	30.29		
1024	cudaMalloc	110.43	94.21		
	cudaLaunchKernel	6.40E+00	5.46		
32768	cudaMemcpy HtoD	2.62E-02	66.21		
	cudaMemcpy DtoH	1.11E-02	27.93		
	saxpy_gpu()	2.34E-03	5.86		
32768	cudaMalloc	92.967	93.32		
	cudaLaunchKernel	6.0815	6.1		
131072	cudaMemcpy HtoD	9.05E-02	66.99		
	cudaMemcpy DtoH	4.12E-02	30.55		
	saxpy_gpu()	3.33E-03	2.46		
131072	cudaMalloc	103.37	92.32		
	cudaLaunchKernel	6.7579	6.04		

*Table 1: SAXPY Raw Data from nvprof Profiling Tool*

This data is further synthesized into total execution time from both Kernel and API calls for each input size (see Table 2). A notable trend is that the GPU execution time scales with input size but the API calls took largely the same time despite scaling the input size. Indeed when profiling the data, the input size for API made no difference in the time it took and varied by ~20ms randomly each program run through.





Table 2: SAXPY Kernel and API Execution Time for Various Input Sizes

The breakdown of top time consumers for Kernel and API for each input size is shown below in Graph 1. The biggest consumers across input size were memory operations. For the API cudaMalloc always took the most time and for the Kernel cudaMemcpy HtoD was the biggest consumer (though to a lesser extent).



Graph 1: SAXPY Execution Time Breakdown by Top Consumers

### Monte-Carlo Execution:

The raw data from profiling the Monte-Carlo program is shown below in Table 3. For three input sizes (sampleSize), the data ranges from 1000 to 1E9. The time execution took increased greatly between 1E6 and 1E9 but not between 1000 and 1E6.

Monte-Carlo Program						
Input Size	Kernel / API Call	Execution Time (ms)	Percentage of Total Time (%)			GPU
1.00E+03	cudaMemcpy DtoH	2.30E-03	1.06			API
1.00E+03	reduceCounts()	8.45E-03	3.9			
1.00E+03	generatePoints()	2.06E-01	95.03			
1.00E+03	cudaMalloc	104.48	93.96			
1.00E+03	cudaMemcpy	6.3282	5.69			
1.00E+06	cudaMemcpy DtoH	3.24E+00	4.88			
1.00E+06	generatePoints()	6.32E+01	95.11			
1.00E+06	cudaMalloc	108.19	59.33			
1.00E+06	cudaMemcpy	67.58	37.06			
1.00E+06	cudaLaunchKernel	6.2055	3.4			
1.00E+09	cudaMemcpy DtoH	2.67E+03	4.52			
1.00E+09	generatePoints()	5.63E+04	95.46			
1.00E+09	cudaMalloc	109.34	0.19			
1.00E+09	cudaMemcpy	5.89E+04	99.79			

Table 3: Monte-Carlo Raw Data from nvprof Profiling Tool

This gap in size can be shown clearly in Table 4. The smaller two input sizes merged into one data point near the origin, dwarfed by the larger input size and it's much longer execution time. Interestingly, this program shows that varying sampleSize in the code affects both API and Kernel unlike the SAXPY program which largely only affected Kernel calls.





Table 4: Monte-Carlo Kernel and API Execution Time for Various Input Sizes

The breakdown of top time consumers for Kernel and API for each input size in the Monte-Carlo program is shown below in Graph 2. The biggest consumers across input size were not always memory operations unlike SAXPY. For the API `cudaMalloc` started out as the biggest consumer of time but then steadily decreases with input size and `cudaMemcpy` took over. For the Kernel `generatePoints()` was the biggest consumer every time.

	<code>cudaMemcpy DtoH</code>	<code>reduceCounts()</code>	<code>generatePoints()</code>	<code>cudaMalloc</code>	<code>cudaMemcpy</code>	<code>cudaLaunchKernel</code>
1000 Kernel	1.06	3.9	95.03	0	0	0
1000 API	0	0	0	93.96	5.69	0
1000000 Kernel	4.88	0	95.11	0	0	0
1000000 API	0	0	0	59.33	37.06	3.4
1000000000 Kernel	4.52	0	95.46	0	0	0
1000000000 API	0	0	0	0.19	99.79	0



*Graph 2: Monte-Carlo Execution Time Breakdown by Top Consumers*